



# Reproduce of ESync

**Paper:** ESync: Accelerating Intra-Domain Federated Learning in Heterogeneous Data Centers

**Team:** Li Kang(models&trainers), Zhao Tian.Y(state server&deployment), Lee Sen.J(state server&architecture)

**Date:** 31th / October / 2024



# Outline

1. Look Back at the Paper
2. Architecture Design
3. Algorithm Implementation
4. Experimental Evaluation
5. Appreciation



# Background of the paper

## Scenario and Feature

- Cross-device FL & Cross-Silo FL
  - Bandwidth-limited **inter-domain** networks (cross WAN),
  - Isolated parties (scale at hundreds or thousands),
  - Uneven and non-i.i.d data distribution.
- Novel scenarios: **intra-domain FL** in heterogeneous data centers
  - Sufficient bandwidth,
  - Highly heterogeneous computing power.

# Idea

Combining the ideas of existing either synchronous or asynchronous algorithms, the paper proposes the **ESync** algorithm.




Fig. 7. ESync timeline with adaptive local iterations.

# Visualization of the idea




Fig. 6. Illustration of (a) SSGD and (b) ESync in the early training, and (c) ESync in the later training.

## How to implement ESync?

- A novel **scheduler**.
- An efficient **synchronization algorithm** (based on SSGD).
- Analyze the **trade-off** between convergence accuracy and communication efficiency.

# Architecture Design

## Overview

- **State Server:** monitor the status and progress, coordinate the synchronization pace of all workers.
- **Parameter Server:** average the updates and synchronize the averaged updates to all workers.
- **Workers:** perform local updates and send their updates to the parameter server.




Fig. 2. Overview of the proposed architecture. ① The worker trains a local model on local data. ② The worker queries the next action from the state server. ③ The state server responds with the next action. ④ The worker pushes an update to the parameter server. ⑤ The worker pulls the averaged update from the parameter server.

# State Server

State Server contains

- a **message receiver**, a **message sender**,
- a **FIFO message queue**,
  - buffer the received messages.
- a **message router**,
  - forward to the handlers according to `msg_type` field.
- a **state database**,
  - lock-free state table.
- and a set of **message handlers**.

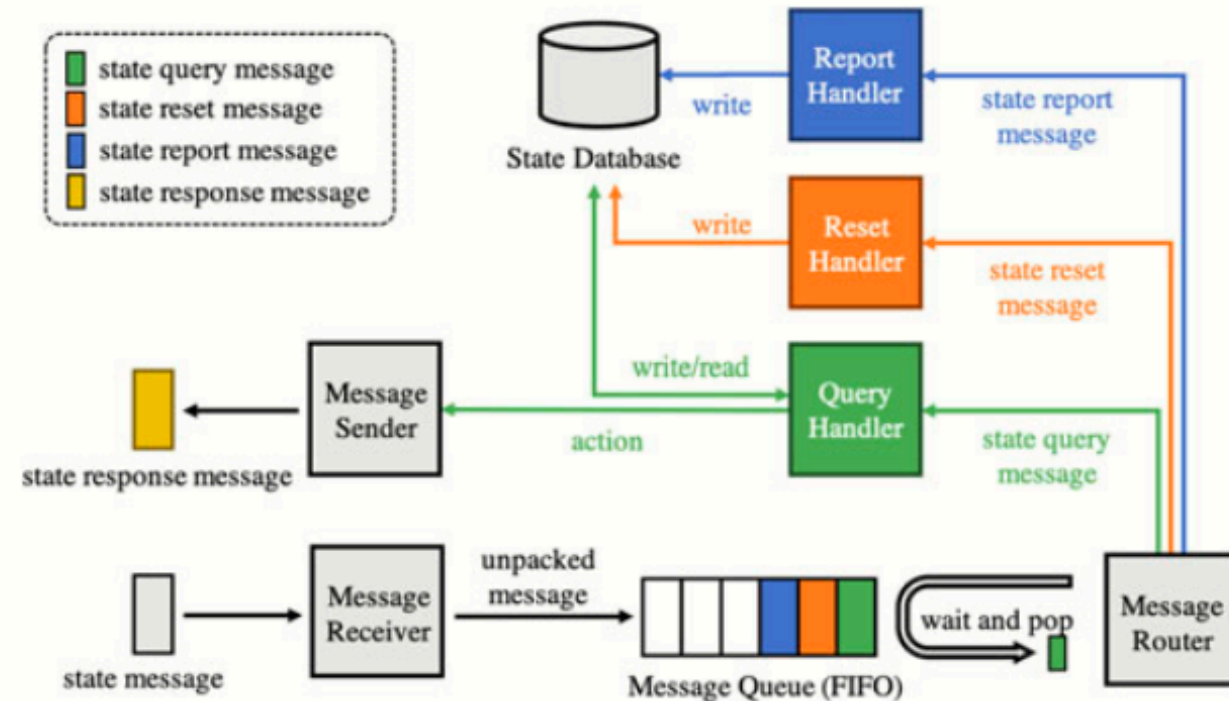


Fig. 3. Implementation of state server.

# Message Format & Message Types

- **state reset message**: initiated by the Parameter Server, to initialize the records in the state database.
- **~ report msg**: initiated by the worker to synchronize its status and progress to State Server.
- **~ query msg**: initiated by the worker to query State Server for the next action.
- **~ response msg**: used by State Server to inform the querying worker of the next action.

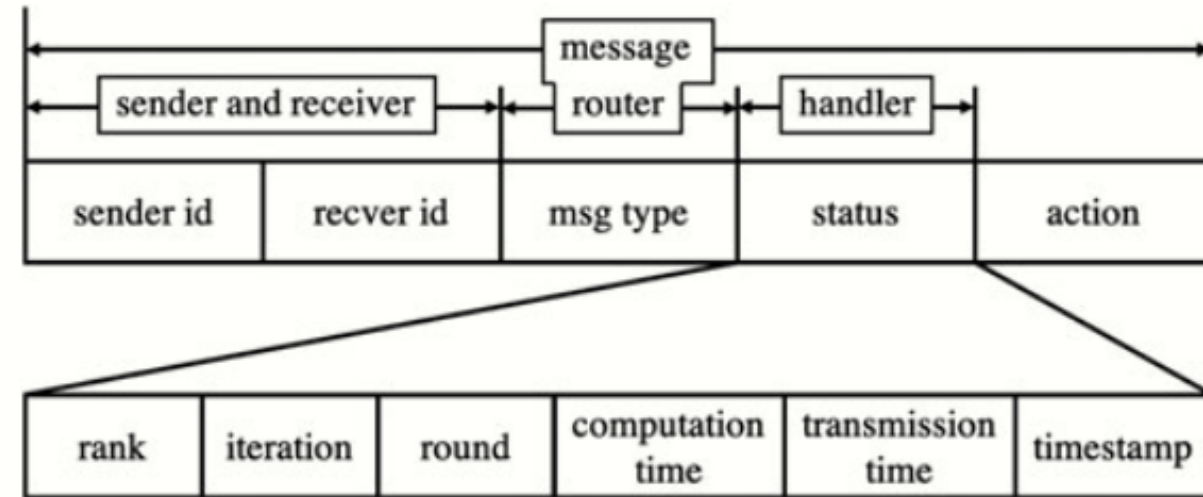


Fig. 4. The structure of message.





# Algorithm Implementation

## Project Overview

### ESync-Reproduce

- |— datasets/
- |— images/
- |— logs/
- |— StateServer/
- |— models/
- |— trainer/
- |— main.py, config.py, draw.py, summary.py, utils.py

# Dataset Selection

## ESync-Reproduce

```
|— data/  
    |— FashionMNIST/, CIFAR10/, ...  
|— output/  
|— logs/  
|— StateServer/  
|— models/  
|— trainer/  
|— main.py, config.py, draw.py, summary.py, utils.py
```

We chose to test the performance of the **ESync** algorithm in image classification training. So, we use the FashionMNIST, CIFAR10 ,etc. datasets for experiments.



# StateServer Implementation

```
|—— StateServer/  
    |—— __init__.py  
    |—— settings.py  
    |—— urls.py # Django URL configuration file.  
    |—— wsgi.py # Django Entry Point.  
    |—— views.py # Core Functionality.  
    |—— manage.py
```

In the `views.py`, we implement the core functionality of the StateServer.

- Global variables: `num_workers`, `records`, `t`, `epsilon`.
- Functions: `init_state_server(request)`, `find_stragglers()`, `apply_for_aggregation(request)`, `reset_state_server(request)`,



# Models Implementation

```
|— models/  
    |— alexnet.py  
    |— resnet.py  
    |— ...
```

Due to the limitation of time, we conducted experiments using only `alexnet` and `resnet` from existing frameworks, adapted for training in federated learning with these frameworks.

Hyperparameters for `resnet`:

- learning rate: 0.001, local lr: 0.001, global lr: 1.0,
- batch size: 64, evaluation duration: 1,
- state server ip: 172.17.73.161, state server port: 9091.



# Trainer Implementation

```
|— trainer/  
    |— esync_trainer.py  
    |— async_trainer.py  
    |— sync_trainer.py  
    |— local_trainer.py
```

Taking `esync_trainer.py` as example, it contains the implementation of a training process for a machine learning model in a distributed environment.

- Imports: `time`, `requests`, `mxnet`, `utils`.
- Variables: `local_lr`, `global_lr`, `batch_size`, etc.
- Training Loop: Batch Processing, Optimization, Global Synchronization.
- State Report, Response Handling, Evaluation(accuracy, loss, time).



# Distributed Training

## 1. Start the StateServer.

```
python3 SimpleStateServer/manage.py runserver 0.0.0.0:10010
```

The state server will listen on port 10010 in the background to wait for the queries from workers.

## 2. Start the Schedulers.

```
> DMLC_ROLE=scheduler DMLC_PS_ROOT_URI=172.17.73.161 DMLC_PS_ROOT_PORT=9091 DMLC_NUM_SERVER=1 \  
DMLC_NUM_WORKER=4 PS_VERBOSE=1 DMLC_INTERFACE=eth0 \  
python3 main.py > scheduler.log &
```

Start the scheduler and listen on port 9091 to wait for the messages from workers.



### 3. Start the Servers.

```
> DMLC_ROLE=server DMLC_PS_ROOT_URI=172.17.73.161 DMLC_PS_ROOT_PORT=9091 DMLC_NUM_SERVER=1 \
DMLC_NUM_WORKER=4 PS_VERBOSE=1 DMLC_INTERFACE=eth0 \
python3 main.py > server.log &
```

The parameter `server` will create a TCP connection to the scheduler and complete registration automatically by specifying **DMLC\_PS\_ROOT\_URI** and **DMLC\_PS\_ROOT\_PORT** (the same as workers). The aggregation operations will be performed on GPU if **cpu** is set to *False*.

### 4. Start the Workers.

```
> DMLC_ROLE=worker DMLC_PS_ROOT_URI=172.17.73.161 DMLC_PS_ROOT_PORT=9091 DMLC_NUM_SERVER=1 \
DMLC_NUM_WORKER=4 PS_VERBOSE=1 DMLC_INTERFACE=eth0 \
python3 main.py -m esync -g 0 -n resnet18-v1 -ll 0.001 -b 64 -dcasgd 0 -s 0
```



# Dockerization

We provide a Dockerfile to build the image for the StateServer and the worker.

```
FROM python:3.6
COPY . .
RUN pip install -r requirements.txt
ENTRYPOINT ["python", "main.py", "-dd", "./data/"]
CMD ["-m", "local", "-c", "1", "-e", "1000"]
```

```
tonyzhao@rupa:~/ESync-master$ sudo docker build . -t esync
[sudo] password for tonyzhao:
[+] Building 29.7s (8/8) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 271B
=> [internal] load metadata for dockerproxy.net/library/python:3.6
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load build context
=> => transferring context: 2.58kB
=> CACHED [1/3] FROM dockerproxy.net/library/python:3.6@sha256:f8652afaf88c25f0d22354d547d892591067aa4026a7fa9a6819df9f300af6fc
=> [2/3] COPY . .
=> [3/3] RUN pip install -r requirements.txt -i https://mirrors.sustech.edu.cn/pypi/web/simple
=> exporting to image
=> => exporting layers
=> => writing image sha256:0426194aed524b82aadca7c036f2e11951d1e8c0514556ed04d32008b2d079aa
=> => naming to dockerproxy.net/library/esync
```



# Experimental Evaluation

## Experimental Setup

- **Environment Setup:**
- **Models and Datasets:** AlexNet, ResNet, CIFAR10, FashionMNist.
- **Benchmark Algorithms:** SSGD, ASGD.
- **Training Hyperparameters:** ...


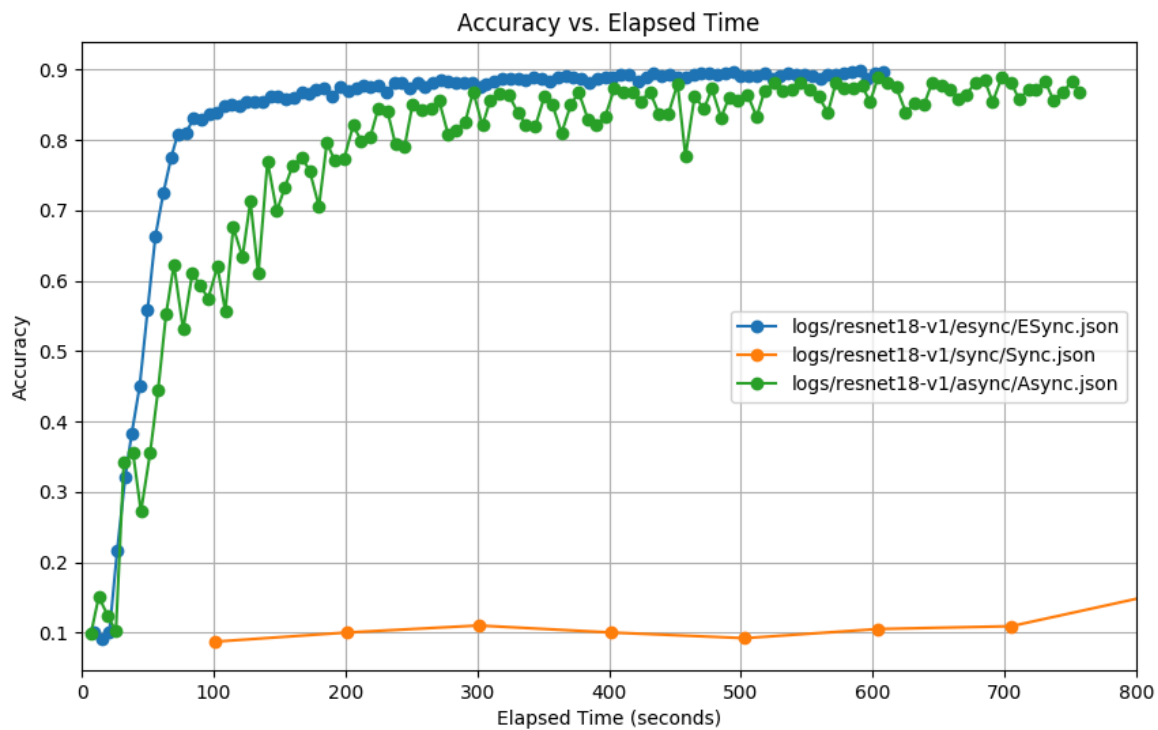
### 配置

4核(vCPU) 8 GiB 5 Mbps  
[ecs.e-c1m2.xlarge](#)

2核(vCPU) 8 GiB 5 Mbps  
GPU: NVIDIA P4  
[ecs.gn5i-c2g1.large](#)

2核(vCPU) 8 GiB 5 Mbps  
GPU: NVIDIA P4  
[ecs.gn5i-c2g1.large](#)

# Numerical Results





**Thanks for your attention!**